

In [2]:

```
pwd
```

Out[2]:

```
'C:\\Users\\Acer'
```

In [3]:

```
! pip install xgboost
```

```
Requirement already satisfied: xgboost in d:\\python\\lib\\site-packages (1.5.2)
```

```
Requirement already satisfied: numpy in d:\\python\\lib\\site-packages (from xgboost) (1.19.5)
```

```
Requirement already satisfied: scipy in d:\\python\\lib\\site-packages (from xgboost) (1.5.2)
```

In [4]:

```
! pip install dtreeviz
```

```
Requirement already satisfied: dtreeviz in d:\python\lib\site-packages (1.3.3)
Requirement already satisfied: matplotlib in d:\python\lib\site-packages (from dtreeviz) (3.3.2)
Requirement already satisfied: graphviz>=0.9 in d:\python\lib\site-packages (from dtreeviz) (0.19.1)
Requirement already satisfied: pytest in d:\python\lib\site-packages (from dtreeviz) (0.0.0)
Requirement already satisfied: scikit-learn in d:\python\lib\site-packages (from dtreeviz) (0.24.2)
Requirement already satisfied: pandas in d:\python\lib\site-packages (from dtreeviz) (1.3.5)
Requirement already satisfied: numpy in d:\python\lib\site-packages (from dtreeviz) (1.19.5)
Requirement already satisfied: colour in d:\python\lib\site-packages (from dtreeviz) (0.1.5)
Requirement already satisfied: certifi>=2020.06.20 in d:\python\lib\site-packages (from matplotlib->dtreeviz) (2021.5.30)
Requirement already satisfied: python-dateutil>=2.1 in c:\users\acer\appdata\roaming\python\python38\site-packages (from matplotlib->dtreeviz) (2.8.2)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in d:\python\lib\site-packages (from matplotlib->dtreeviz) (2.4.7)
Requirement already satisfied: kiwisolver>=1.0.1 in d:\python\lib\site-packages (from matplotlib->dtreeviz) (1.3.0)
Requirement already satisfied: pillow>=6.2.0 in d:\python\lib\site-packages (from matplotlib->dtreeviz) (8.0.1)
Requirement already satisfied: cyclert>=0.10 in d:\python\lib\site-packages (from matplotlib->dtreeviz) (0.10.0)
Requirement already satisfied: attrs>=17.4.0 in d:\python\lib\site-packages (from pytest->dtreeviz) (20.3.0)
Requirement already satisfied: iniconfig in d:\python\lib\site-packages (from pytest->dtreeviz) (1.1.1)
Requirement already satisfied: packaging in d:\python\lib\site-packages (from pytest->dtreeviz) (20.4)
Requirement already satisfied: pluggy<1.0,>=0.12 in d:\python\lib\site-packages (from pytest->dtreeviz) (0.13.1)
Requirement already satisfied: py>=1.8.2 in d:\python\lib\site-packages (from pytest->dtreeviz) (1.9.0)
Requirement already satisfied: toml in c:\users\acer\appdata\roaming\python\python38\site-packages (from pytest->dtreeviz) (0.10.2)
Requirement already satisfied: atomicwrites>=1.0 in d:\python\lib\site-packages (from pytest->dtreeviz) (1.4.0)
Requirement already satisfied: colorama in c:\users\acer\appdata\roaming\python\python38\site-packages (from pytest->dtreeviz) (0.4.4)
Requirement already satisfied: threadpoolctl>=2.0.0 in d:\python\lib\site-packages (from scikit-learn->dtreeviz) (2.1.0)
Requirement already satisfied: scipy>=0.19.1 in d:\python\lib\site-packages (from scikit-learn->dtreeviz) (1.5.2)
Requirement already satisfied: joblib>=0.11 in d:\python\lib\site-packages (from scikit-learn->dtreeviz) (0.17.0)
Requirement already satisfied: pytz>=2017.3 in d:\python\lib\site-packages (from pandas->dtreeviz) (2020.1)
Requirement already satisfied: six>=1.5 in c:\users\acer\appdata\roaming\python\python38\site-packages (from python-dateutil>=2.1->matplotlib->dtreeviz) (1.16.0)
```

In [5]:

```
from dtreeviz.trees import dtreeviz
```

In [6]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.ticker as mtick
import matplotlib.pyplot as plt
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from imblearn.over_sampling import SMOTE
from sklearn.metrics import fbeta_score
from numpy import loadtxt
from xgboost import XGBClassifier
```

In [7]:

```
telco_base_data = pd.read_csv('D:/Customer_churn_analysis/WA_Fn-UseC_-Telco-Customer-Churn.
```

In [8]:

```
telco_base_data.head()
```

Out[8]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	Nc
2	3668-QPYBK	Male	0	No	No	2	Yes	Nc
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	Nc

5 rows × 21 columns



In [9]:

```
telco_base_data.shape
```

Out[9]:

(7043, 21)

In [10]:

```
telco_base_data.columns.values
```

Out[10]:

```
array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',  
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',  
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',  
      'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
      'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',  
      'TotalCharges', 'Churn'], dtype=object)
```

describe is used for numerical data

In [11]:

```
telco_base_data.dtypes
```

Out[11]:

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int64
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object
StreamingTV	object
StreamingMovies	object
Contract	object
PaperlessBilling	object
PaymentMethod	object
MonthlyCharges	float64
TotalCharges	object
Churn	object
dtype:	object

In [12]:

```
telco_base_data.describe()
```

Out[12]:


	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

In [13]:

```
telco_base_data['Churn'].value_counts().plot(kind='barh', figsize=(8,6))  
plt.xlabel("Count",labelpad=14)  
plt.ylabel("Target Variable",labelpad=14)  
plt.title("Count of target variable",y=1.02)
```

Out[13]:

```
Text(0.5, 1.02, 'Count of target variable')
```



In [14]:

```
telco_base_data['Churn'].value_counts()
```

Out[14]:

```
No      5174
Yes     1869
Name: Churn, dtype: int64
```

In [15]:

```
100*telco_base_data['Churn'].value_counts()/len(telco_base_data['Churn'])
# Basically value counts is used to return unique values. So there are only 2 unique values
```

Out[15]:

```
No      73.463013
Yes     26.536987
Name: Churn, dtype: float64
```

imbalanced dataset

In [16]:

```
telco_base_data.info(verbose = True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customerID            7043 non-null   object
 1   gender                7043 non-null   object
 2   SeniorCitizen         7043 non-null   int64
 3   Partner               7043 non-null   object
 4   Dependents            7043 non-null   object
 5   tenure               7043 non-null   int64
 6   PhoneService          7043 non-null   object
 7   MultipleLines         7043 non-null   object
 8   InternetService       7043 non-null   object
 9   OnlineSecurity        7043 non-null   object
10   OnlineBackup          7043 non-null   object
11   DeviceProtection      7043 non-null   object
12   TechSupport           7043 non-null   object
13   StreamingTV           7043 non-null   object
14   StreamingMovies       7043 non-null   object
15   Contract              7043 non-null   object
16   PaperlessBilling      7043 non-null   object
17   PaymentMethod         7043 non-null   object
18   MonthlyCharges        7043 non-null   float64
19   TotalCharges          7043 non-null   object
20   Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

In [17]:

```
telco_data =telco_base_data.copy()
```

In [18]:

```
telco_data.TotalCharges = pd.to_numeric(telco_data.TotalCharges, errors='coerce')
telco_data.isnull().sum()
```

Out[18]:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0

dtype: int64

In [19]:

```
telco_data.loc[telco_data ['TotalCharges'].isnull()== True]
```

Out[19]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLi
488	4472-LVYGI	Female	0	Yes	Yes	0	No	No pr ser
753	3115-CZMZD	Male	0	No	Yes	0	Yes	
936	5709-LVOEQ	Female	0	Yes	Yes	0	Yes	
1082	4367-NUYAO	Male	0	Yes	Yes	0	Yes	
1340	1371-DWPAZ	Female	0	Yes	Yes	0	No	No pr ser
3331	7644-OMVMY	Male	0	Yes	Yes	0	Yes	
3826	3213-VVOLG	Male	0	Yes	Yes	0	Yes	
4380	2520-SGTTA	Female	0	Yes	Yes	0	Yes	
5218	2923-ARZLG	Male	0	Yes	Yes	0	Yes	
6670	4075-WKNIU	Female	0	Yes	Yes	0	Yes	
6754	2775-SEFEE	Male	0	No	Yes	0	Yes	

11 rows × 21 columns

since the nuber 11 is quite negligible as compared to 7413 we can ignore them

In [20]:

```
telco_data.dropna(how = 'any', inplace = True)
# telco_data.fillna(0)
```

In [21]:

```
#creating Bins
```

In [22]:

```
print(telco_data['tenure'].max())
```

72

In [23]:

```
labels = ["{0} - {1}".format(i,i+11) for i in range(1,72,12)]
telco_data['tenure_group'] = pd.cut(telco_data.tenure, range(1,80,12), right=False, labels=labels)
```

In [24]:

```
telco_data['tenure_group'].value_counts()
```

Out[24]:

```
1 - 12      2175
61 - 72     1407
13 - 24     1024
25 - 36      832
49 - 60      832
37 - 48      762
Name: tenure_group, dtype: int64
```

In [25]:




```
telco_data.drop(columns= ['customerID','tenure'], axis=1, inplace=True)
telco_data.head()
```




Out[25]:




	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	On
0	Female	0	Yes	No	No	No phone service	DSL	
1	Male	0	No	No	Yes	No	DSL	
2	Male	0	No	No	Yes	No	DSL	
3	Male	0	No	No	No	No phone service	DSL	
4	Female	0	No	No	Yes	No	Fiber optic	




In [26]:





```
for i, predictor in enumerate(telco_data.drop(columns=['Churn', 'TotalCharges', 'MonthlyCharg  
plt.figure(i)  
sns.countplot(data=telco_data, x=predictor, hue='Churn')
```














convert The variable Churn in a binary numeric Value: Yes=1, No=0

In [27]:

```
telco_data['Churn'] = np.where(telco_data.Churn== 'Yes',1,0)
```

In [28]:

```
telco_data.head()
```

Out[28]:

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	On
0	Female	0	Yes	No	No	No phone service	DSL	
1	Male	0	No	No	Yes	No	DSL	
2	Male	0	No	No	Yes	No	DSL	
3	Male	0	No	No	No	No phone service	DSL	
4	Female	0	No	No	Yes	No	Fiber optic	

convert all the categorical vaiables into numerical/dummy variables

In [29]:

```
#one hot Encoding
```

In [30]:

```
#dummy Trap can be used
```


In [31]:

```
telco_data_dummies = pd.get_dummies(telco_data)
telco_data_dummies.head()
```

Out[31]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_M
0	0	29.85	29.85	0	1	0	
1	0	56.95	1889.50	0	0	1	
2	0	53.85	108.15	1	0	1	
3	0	42.30	1840.75	0	0	1	
4	0	70.70	151.65	1	1	0	

5 rows × 51 columns




In [32]:

```
sns.lmplot(data=telco_data_dummies, x='MonthlyCharges', y='TotalCharges', fit_reg=False)
```

Out[32]:

<seaborn.axisgrid.FacetGrid at 0x17fc1c2cf40>




In [33]:

```
math = sns.kdeplot(telco_data_dummies.MonthlyCharges[(telco_data_dummies["Churn"]==0)],  
                  color="Red", shade = True)  
math = sns.kdeplot(telco_data_dummies.MonthlyCharges[(telco_data_dummies["Churn"]==1)],  
                  ax=math,color="Green", shade = True)  
math.legend(["No churn", "Churn"], loc='upper right')  
math.set_ylabel('Density')  
math.set_xlabel('Monthly Charges')  
math.set_title('Monthly charges by churn')
```

Out[33]:

Text(0.5, 1.0, 'Monthly charges by churn')



In [34]:


```
# key Insight found As higher the monthly charge, Lower the tenure, Lower the Total charge-
```

In [35]:

```
Tot = sns.kdeplot(telco_data_dummies.TotalCharges[(telco_data_dummies["Churn"] == 0) ],  
                  color="Red", shade = True)  
Tot = sns.kdeplot(telco_data_dummies.TotalCharges[(telco_data_dummies["Churn"] == 1) ],  
                  ax =Tot, color="Blue", shade= True)  
Tot.legend(["No Churn", "Churn"],loc='upper right')  
Tot.set_ylabel('Density')  
Tot.set_xlabel('Total Charges')  
Tot.set_title('Total charges by churn')
```

Out[35]:

Text(0.5, 1.0, 'Total charges by churn')



In []:


```
# So from here we can Infer that that There is more churning when the total Charges are Low  
# understanding because that is quite obvious that if you have not paid a hefty amount to s  
# switch.
```

In [36]:

```
plt.figure(figsize=(20,8))
telco_data_dummies.corr()['Churn'].sort_values(ascending = False).plot(kind = 'bar')
```

Out[36]:

<AxesSubplot:>



In [37]:

key Insights:-


High Churn Rate seen in Contract_Month_to_month, No online Security , No Tech Support, Fi
 # Low Churn Rate seen in Long term Contracts, Subscriptions without net, Customer Enged for
 # Gender, Availablity of phone Service,-----> no impact on churn

In [38]:

```
plt.figure(figsize=(12,12))
sns.heatmap(telco_data_dummies.corr(), cmap="Paired")
```

Out[38]:

<AxesSubplot:>



In [39]:

```
new_df1_target0 = telco_data.loc[telco_data["Churn"]==0]
new_df1_target1 = telco_data.loc[telco_data["Churn"]==1]
```

In [40]:


```
def uniplot(df,col,title,hue = None):

    sns.set_style('whitegrid')
    sns.set_context('talk')
    plt.rcParams["axes.labelsize"]= 20
    plt.rcParams['axes.titlesize'] = 22
    plt.rcParams['axes.titlepad']= 30

    temp = pd.Series(data=hue)
    fig,ax= plt.subplots()
    width = len(df[col].unique()) +7 + 4*len(temp.unique())
    fig.set_size_inches(width , 8)
    plt.xticks(rotation=45)
    plt.yscale('log')
    plt.title(title)
    ax = sns.countplot(data = df, x= col, order= df[col].value_counts().index,hue = hue, pa
    plt.show()
```

In [41]:

```
uniplot(new_df1_target1,col='Partner',title='Distribution of Gender for Churned Customers',
```




In [42]:

```
# So from here we can see that if there is no partner and and if the selected candidate is
# the Inverse.
```


In [43]:

```
unipLOT(new_df1_target1,col='PaymentMethod',title= 'Distribution of PaymentMethod for Churn
```




In [44]:

```
unipLOT(new_df1_target1,col='Contract',title= 'Distribution of Contract for Churned Custome
```



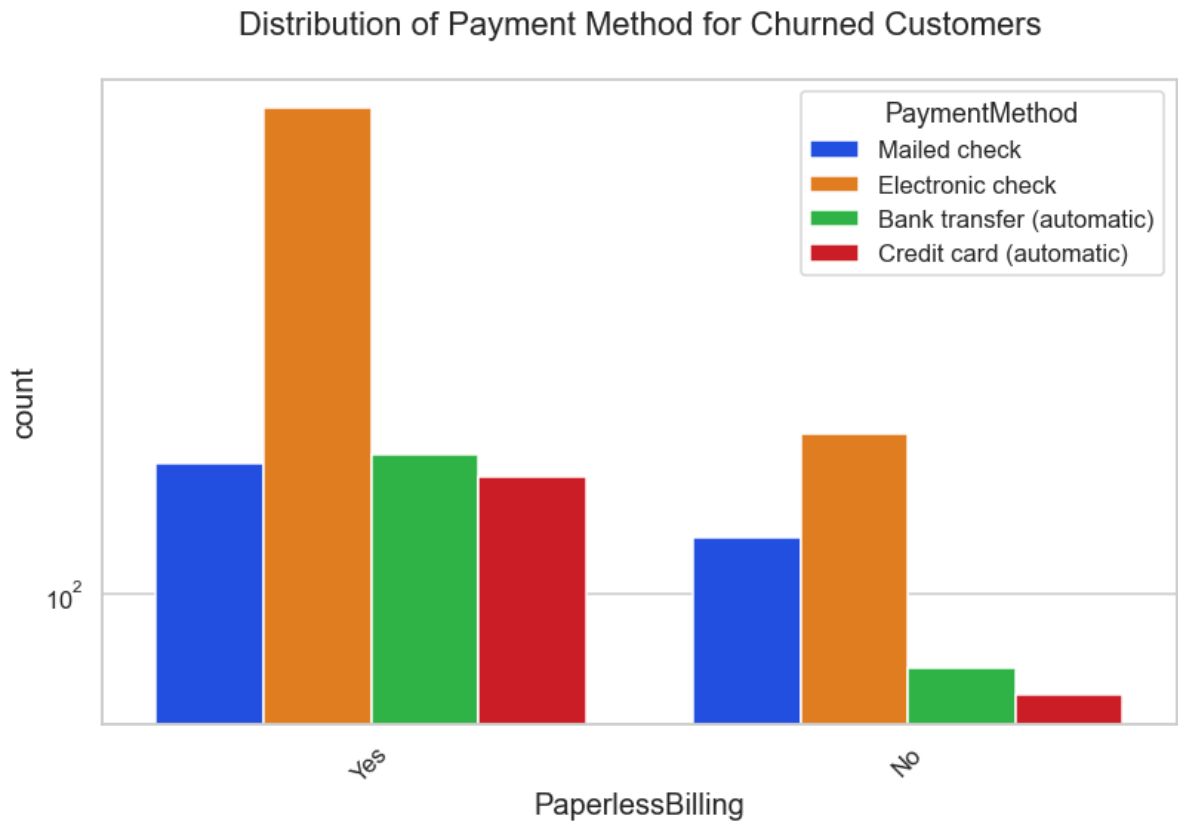
In [45]:

```
unipLOT(new_df1_target1,col='SeniorCitizen',title= 'Distribution of SeniorCitizen for Churn
```




In [46]:

```
unipLOT(new_df1_target1,col='PaperlessBilling',title= 'Distribution of Payment Method for C
```



In [47]:

```
unipLOT(new_df1_target1,col='StreamingMovies',title= 'Distribution of Internet Service for C
```



In [48]:

```
# Basic Conclusions from Univariate and Bivariate analysis  
# 1.) Monthly Contract People Are most likely to Churn at higher rate  
# 2.) Electronic Check Payment are the People having Higher Churning rate  
# 3.) Non- Senior Citizen Citizen are High Churners  
# 4.) No Online Security , No Tech Support category are higher Churners
```

In [49]:

```
telco_data_dummies.to_csv('tel_churn.csv')
```

In [50]:

```
df = pd.read_csv('tel_churn.csv')
```


In [51]:

```
df.head()
```

Out[51]:

	Unnamed: 0	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male
0	0	0	29.85	29.85	0	1	0
1	1	0	56.95	1889.50	0	0	1
2	2	0	53.85	108.15	1	0	1
3	3	0	42.30	1840.75	0	0	1
4	4	0	70.70	151.65	1	1	0

5 rows × 52 columns

In [52]:

```
df = df.drop('Unnamed: 0', axis=1)
```

In [53]:

```
df.head()
```

Out[53]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_M
0	0	29.85	29.85	0	1	0	
1	0	56.95	1889.50	0	0	1	
2	0	53.85	108.15	1	0	1	
3	0	42.30	1840.75	0	0	1	
4	0	70.70	151.65	1	1	0	

5 rows × 51 columns

In [54]:

```
# now create variables x and y where x is independent variables and y is dependent variable
```


In [55]:

```
x= df.drop('Churn', axis=1)
x
```

Out[55]:

	SeniorCitizen	MonthlyCharges	TotalCharges	gender_Female	gender_Male	Partner_No	F
0	0	29.85	29.85	1	0	0	
1	0	56.95	1889.50	0	1	1	
2	0	53.85	108.15	0	1	1	
3	0	42.30	1840.75	0	1	1	
4	0	70.70	151.65	1	0	1	
...
7027	0	84.80	1990.50	0	1	0	
7028	0	103.20	7362.90	1	0	0	
7029	0	29.60	346.45	1	0	0	
7030	1	74.40	306.60	0	1	0	
7031	0	105.65	6844.50	0	1	1	

7032 rows × 50 columns



In [56]:

```
y = df['Churn']
y
```

Out[56]:

```
0      0
1      0
2      1
3      0
4      1
..
7027   0
7028   0
7029   0
7030   1
7031   0
```

Name: Churn, Length: 7032, dtype: int64

In [57]:

```
# Evaluation Metrics to be considered
```

In [58]:

```
# Before Moving on to the model processing we must be sure that on what basis are we going
# If we observe our output column is Churn that is a categorical column so we sevrsl choice
# Recall , F-1/F-2/F-0.5 scores.
```

In [59]:

```
# Lets try to understand by considering the first evaluation metric that is accuracy.
# Accuracy will not be a good evaluation metric the reason being the Imbalanced dataset. If
# model will be more biased towards the majority class that is not churning and it is high
# minority class. Due to which even though we will have high accuracy but that would not ad
# won't be a deciding parameter.
```

In [60]:

```
# Let us try with the second Evaluation metric that is Precision.
# Precision means out of all the predicted positive values how many of them are actually po
# So if we try to observe the Formula of Precision :-  $T.P / T.P + F.P$ 
# So we are concerned more with the term False positive :- Here False positive means we pre
# but actually he is not. So if we consider this from a business point of view then there i
# though our model is bit inaccurate.
# so we will not consider Precision in this case.
```

In [61]:

```
# Let us try with the recall :
# Recall means out of actual positive values how many of them are correctly classified.
# So if we look at the formula of Recall:-  $T.P / T.P + F.N$ 
# Here we are concerned more with the false negative term :- False negative means we predic
# but actually he is churning. So this could be a major Setback for the organization becaus
# Loyal customer where as he is a churner.
# So this is a potential parameter that could be considered as evaluation metric
```

In [62]:

```
# Now we know that Recall could be a Important evaluation metric and in our scenario its mo
# So we are not neglecting the F- score because the number of False positive is not that hi
# So if we are considering the F-score as  $F.N > F.P$  (Weightage) So we will be considering F-2,
# balance.
```

In [63]:

```
# Normal train test Split (70-30)
```

In [64]:

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3)
```

In []:

```
# So as we see that recall score is quite less (less than 50%) so the amount of False negat
# our models recall score is less. Due to which it becomes quite Important to emphasize on
# business perspective of an organization in a wrong direction.
```

In [105]:

```
# Decision Tree Classifier
```

In [106]:

```
model_dt = DecisionTreeClassifier(criterion='entropy', random_state=100, max_depth=6, min_sam
```

In [107]:

```
model_dt.fit(x_train, y_train)
```

Out[107]:

```
DecisionTreeClassifier(criterion='entropy', max_depth=6, min_samples_leaf=8,
                      random_state=100)
```

In [108]:

```
y_pred = model_dt.predict(x_test)
```

In [109]:

```
y_pred
```

Out[109]:

```
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

In [110]:

```
model_dt.score(x_test, y_test)
```

Out[110]:

```
0.7895734597156399
```

In [111]:

```
print(classification_report(y_test, y_pred, labels = [0,1]))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1557
1	0.63	0.49	0.55	553
accuracy			0.79	2110
macro avg	0.73	0.69	0.71	2110
weighted avg	0.78	0.79	0.78	2110

In [115]:

```
confusion=confusion_matrix(y_test, y_pred)
print(confusion)
```

```
[[1395 162]
 [ 282 271]]
```

In [117]:

```
recall = confusion[0][0] / (confusion[0][0] + confusion[1][0])
print(recall)
```

```
0.8318425760286225
```

In []:

```
# So we obtained the Recall for the Decision Tree Model
```

In []:

```
# Recall value should Ideally be high because it tells that the there are less false negati
```

In []:

```
# 80% recall means that our model has correctly classified 80% of actual positive data. To o
# Differnt models and then can test the result.
```

In [118]:

```
#RANDOM FOREST CLASSIFIER
```

In [119]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [120]:

```
model_rf = RandomForestClassifier(criterion='gini',random_state=100,max_depth=5, min_sample
model_rf.fit(x_train, y_train)
y_pred_rf = model_rf.predict(x_test)
```

In [121]:

```
model_rf.score(x_test,y_test)
```

Out[121]:

```
0.7862559241706161
```

In [122]:

```
print(classification_report(y_test, y_pred, labels = [0,1]))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1557
1	0.63	0.49	0.55	553
accuracy			0.79	2110
macro avg	0.73	0.69	0.71	2110
weighted avg	0.78	0.79	0.78	2110

In [124]:

```
conf_rf= confusion_matrix(y_test, y_pred_rf)
print(conf_rf)
```

```
[[1472  85]
 [ 366 187]]
```

In [126]:

```
recall = conf_rf[0][0] / (conf_rf[0][0] + conf_rf[1][0])  
print(recall)
```

0.8008705114254625

In [127]:

```
rf = RandomForestClassifier(random_state=42, n_jobs=-1)
```

In [128]:

```
params = {  
    'max_depth': [2,3,5,6,7,15,20],  
    'min_samples_leaf': [5,10,20,50,100,200],  
    'n_estimators': [10,25,30,50,100,200]  
}
```

In [129]:

```
from sklearn.model_selection import GridSearchCV
```

In [130]:

```
grid_search = GridSearchCV(estimator=rf,  
                           param_grid=params,  
                           cv = 4,  
                           n_jobs=-1, verbose=1, scoring="recall")
```

In [131]:

```
grid_search.fit(x_train, y_train)
```

Fitting 4 folds for each of 252 candidates, totalling 1008 fits

Out[131]:

```
GridSearchCV(cv=4, estimator=RandomForestClassifier(n_jobs=-1, random_state=  
42),  
             n_jobs=-1,  
             param_grid={'max_depth': [2, 3, 5, 6, 7, 15, 20],  
                         'min_samples_leaf': [5, 10, 20, 50, 100, 200],  
                         'n_estimators': [10, 25, 30, 50, 100, 200]},  
             scoring='recall', verbose=1)
```

In [133]:

```
rf_best = grid_search.best_estimator_  
rf_best
```

Out[133]:

```
RandomForestClassifier(max_depth=15, min_samples_leaf=10, n_estimators=10,  
                       n_jobs=-1, random_state=42)
```

In [134]:

```
# So we have obtained some of the random forest parameters that will be ideal fit according  
# Lets try with that and see if any change is possible.
```

In [138]:

```
model_rf_1 = RandomForestClassifier(criterion='entropy',random_state=42,max_depth=15, min_s  
model_rf_1.fit(x_train, y_train)  
y_pred_rf_1 = model_rf_1.predict(x_test)
```

In [139]:

```
conf_rf_1= confusion_matrix(y_test, y_pred_rf_1)  
print(conf_rf_1)
```

```
[[1420  137]  
 [ 294  259]]
```

In [140]:

```
recall = conf_rf_1[0][0] /(conf_rf_1[0][0] + conf_rf_1[1][0])  
print(recall)
```

```
0.8284714119019837
```

In []:

```
# so we can see by Hyperparameter tuning(Grid_Search_CV) my Recall score Increased.
```

In []:

```
# Trying with XG boost Classifier
```

In [141]:

```
model = XGBClassifier()
```

In [142]:

```
model.fit(x_train, y_train)
```

D:\PYTHON\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[13:05:28] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Out[142]:

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
              gamma=0, gpu_id=-1, importance_type=None,
              interaction_constraints='', learning_rate=0.300000012,
              max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,
              monotone_constraints=(), n_estimators=100, n_jobs=12,
              num_parallel_tree=1, predictor='auto', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
              tree_method='exact', validate_parameters=1, verbosity=None)
```

In [145]:

```
y_pred = model.predict(x_test)
```

In [146]:

```
conf_XGB = confusion_matrix(y_test, y_pred)
```

In [147]:

```
print(conf_XGB)
```

```
[[1405  152]
 [ 283  270]]
```

In [148]:

```
recall = conf_XGB[0][0] / (conf_XGB[0][0] + conf_XGB[1][0])
print(recall)
```

```
0.832345971563981
```

In []:

```
# So we can observe that by using XG boost classifier we are getting quite decent Recall.
```


In [149]:

```
# So similarly we can try differnt models and try to understand which model can be a perfec
```

In []:

```
# -----
```